

Assignment 1 - Exploring Data. Due September 27, 11:59pm 2019

EPIB607 - Inferential Statistics^a

^aFall 2020, McGill University

This version was compiled on September 11, 2020

All questions are to be answered in an R Markdown document using the provided template and compiled to a pdf document. You are free to choose any function from any package to complete the assignment. Concise answers will be rewarded. Be brief and to the point. Each question is worth 25 points. Label your graphs appropriately with proper titles and axis labels. Please submit both the compiled pdf report to Crowdmark by September 27, 2019, 11:59pm. You need to save your answers to each question in separate pdf files. You also need to upload your code. See <https://crowdmark.com/help/> for details.

Template

Please use the .Rmd template for Assignment 1 is available on myCourses.

1. (25 points) Immunogenicity of the ChAdOx1 nCoV-19 vaccine against SARS-CoV-2

This questions refers to the Lancet paper *Safety and immunogenicity of the ChAdOx1 nCoV-19 vaccine against SARS-CoV-2: a preliminary report of a phase 1/2, single-blind, randomised controlled trial* by Folegatti et. al (2020) and available in myCourses.

- (2 points) Consider Figure 3 Panel B: What visual cues (or aesthetics) are being used? Briefly describe the main takeaways from the entire Figure 3.
- (3 points) Do you think Figure 3 is a good graphic in terms of conveying its message clearly? Is there anything you would have done differently? Explain.
- (2 points) Consider the data introduced in class which contains immunity levels (Immunoglobulin G (IgG)) from the convalescent group and the vaccine groups post 28 days. Note that the IgG levels in the dataset below are given on the log10 scale. Calculate the median IgG levels (ELISA units) on the log10 scale for each group.

```
path <-  
"http://www.biostat.mcgill.ca/hanley/statbook/immunogenicityChAdOx1.nCoV-19vaccine.txt"  
ds <- read.table(path)  
head(ds)
```

#	RefIndexCategory	IgGResponse.log10.ElisaUnits
# 1	Convalescent	2.56
# 2	Convalescent	2.74
# 3	Convalescent	2.79
# 4	Convalescent	3.32
# 5	Convalescent	3.15
# 6	Convalescent	2.35

```
str(ds)
```

```
# 'data.frame': 307 obs. of 2 variables:
# $ RefIndexCategory : Factor w/ 2 levels "Convalescent",...: 1 1 1 1 1 1 1 1
# $ IgGResponse.log10.ElisaUnits: num 2.56 2.74 2.79 3.32 3.15 2.35 2.72 2.95 2.42 2.64
```

- e) (1 point) Are you able to calculate a correlation of IgG levels between groups? If not, explain why not? If yes, interpret the correlation.
- f) (4 points) From the medians alone, is there enough evidence to conclude that the median IgG levels in the convalescent group are higher than the median IgG levels in the vaccine group (post 28 days)? Explain.
- g) (7 points) Use the Bootstrap to assess if there is enough evidence to suggest that the median IgG levels in the convalescent group are higher than the median IgG levels in the vaccine group (post 28 days). *Hint: resample the data with replacement separately in each group B=1000 times. For each of the B datasets, calculate the median IgG level and take the difference in medians between the two groups. Plot the differences in a histogram and calculate the 2.5 and 97.5 percentiles.*
- h) (6 points) The dataset, shown below and available on myCourses, was extracted (approximately) from Figure 3 Panel A for the ChAdOx1 nCoV-19 (prime) group only. The time column represents the days since vaccination, and igg_response are the IgG levels on the original scale. Create an appropriate figure which shows the immunity levels as a function of time. You are free to choose the plot type; the choice of plot should be guided by the message you are trying to convey. Be sure to label your axes, show units, include a title and choose an appropriate color palette. Briefly interpret the plot.

```
DT <- read.csv("prime_igg_response.csv")
```

```
head(DT)
```

```
#   time igg_response
# 1    0   930.37376
# 2    0   267.80142
# 3    0   241.40290
# 4    0   170.80787
# 5    0    79.79795
# 6    0    67.12348
```

2. (25 points) COVID-19 Cases Comparison Between Counties With and Without Stay-at-Home Orders

This question is based on the JAMA Network Open article *Comparison of Estimated Rates of Coronavirus Disease 2019 (COVID-19) in Border Counties in Iowa Without a Stay-at-Home Order and Border Counties in Illinois With a Stay-at-Home Order* by Lyu and Wehby (2020) and available in myCourses. The county and state level cumulative incidence of cases data is provided in the code below. Note: you need to install the covdata package (which is not on CRAN) using `remotes::install_github("kjhealy/covdata")`.

```
# remotes::install_github("kjhealy/covdata")
library(covdata)
library(dplyr); library(tidyr); library(ggplot2); library(readr)
# get population data from https://covid19.census.gov/datasets/
f <- "https://opendata.arcgis.com/datasets/21843f238cbb46b08615fc53e19e0daf_1.csv"
pop_county <- read_csv(file = f) %>%
  dplyr::rename(fips = GEOID, population = B01001_001E, state = State) %>%
  dplyr::select(state, fips, population)

county_level <- nytcovcounty %>%
  dplyr::left_join(pop_county, by = c("state", "fips")) %>%
  dplyr::mutate(cases.per.10k = cases/population * 1e4) %>%
  dplyr::filter(state %in% c("Iowa", "Illinois")) %>%
  dplyr::group_by(county)

pop_state <- pop_county %>%
  dplyr::group_by(state) %>%
  dplyr::summarise(population = sum(population, na.rm = TRUE))

state_level <- county_level %>%
  dplyr::group_by(state, date) %>%
  dplyr::filter(date >= "2020-03-15") %>%
  dplyr::summarise(cases = sum(cases)) %>%
  dplyr::left_join(pop_state, by = "state") %>%
  dplyr::mutate(cases.per.10k = cases / population * 1e4, state = factor(state),
               time = as.numeric(date - min(date)) + 1)
```

```
head(county_level)
```

```
# # A tibble: 6 x 8
# # Groups:   county [1]
#   date      county state   fips  cases deaths population cases.per.10k
#   <date>     <chr> <chr>   <chr> <dbl>  <dbl>      <dbl>      <dbl>
# 1 2020-01-24 Cook    Illinois 17031     1     0    5223719     0.00191
# 2 2020-01-25 Cook    Illinois 17031     1     0    5223719     0.00191
# 3 2020-01-26 Cook    Illinois 17031     1     0    5223719     0.00191
# 4 2020-01-27 Cook    Illinois 17031     1     0    5223719     0.00191
```

```
# 5 2020-01-28 Cook Illinois 17031 1 0 5223719 0.00191
# 6 2020-01-29 Cook Illinois 17031 1 0 5223719 0.00191
```

```
head(state_level)
```

```
# # A tibble: 6 x 6
# # Groups:   state [1]
#   state    date    cases population cases.per.10k  time
#   <fct>   <date>   <dbl>      <dbl>      <dbl> <dbl>
# 1 Illinois 2020-03-15    94    12821497    0.0733     1
# 2 Illinois 2020-03-16   104    12821497    0.0811     2
# 3 Illinois 2020-03-17   159    12821497    0.124      3
# 4 Illinois 2020-03-18   286    12821497    0.223      4
# 5 Illinois 2020-03-19   420    12821497    0.328      5
# 6 Illinois 2020-03-20   583    12821497    0.455      6
```

- (6 points) Using the county level dataset provided, reproduce Figure 1 of the paper. Does your Figure agree with theirs? Would county level curves have been more appropriate to show instead of the state totals?
- (5 points) Plot the cumulative incidence curves per 10000 people from March 21 until the most recent day for which you have data, for each of the counties used in the paper. Interpret the plot and discuss if the county level plots still agree with the overall conclusion of the paper.
- (4 points) Case counts are inherently tied to testing capacity. Death from COVID19 doesn't have this issue, although there are other biases such as misclassification and under reporting. Plot the same graph as in part (b) but for deaths and interpret the plot.
- (10 points) Illinois (Democrat-controlled legislature) is surrounded by states with **Republican-controlled legislatures (Iowa, Missouri, Kentucky, Indiana, Wisconsin)**. Do the data suggest there is a correlation between COVID-19 cases (or deaths) and which party has legislative control? Explain and justify using summary statistics and/or figures. Do not overcomplicate this analysis, i.e., feel free to make simplifying assumptions about testing.

3. (25 points) Age-structures of Populations, then and now

The 1911 census of Ireland was taken on April 2nd 1911 and was released to the public in 1961. Follow [this link](#) for further details on the census. James Hanley (JH) has scrapped the data for Dublin, collected the age-frequency distribution by gender and provided you with a three column .csv file on myCourses called `age_sex_frequencies_ireland.csv` which looks like this:

```
cens <- read.csv("age_sex_frequencies_ireland.csv")
```

```
head(cens)
```

```
#   Gender Age Freq
# 1   Male  0 5332
# 2   Male  1 4570
# 3   Male  2 4979
# 4   Male  3 4789
# 5   Male  4 4884
# 6   Male  5 4787
```

The Age column represents the age in 1911. The Freq column gives the frequency of the number of people for a given age and Gender. Note that Age is an interval; for example, Age=0 actually represents individuals who are between the ages of 0 and 1, Age=1 are individuals between ages 1 and 2, and so on.

- (6 points) What was the earliest year of birth for (i) males and (ii) females ?
- (8 points) Create a suitable visualization of this data and then comment on any patterns you see and give reasons for these patterns. Your choice should leverage all the information provided in the data and be influenced by the message that you are trying to convey. Be sure to include an informative title and figure caption.
- (8 points) Calculate the mean age, the standard deviation (SD), and the quartiles: Q_{25} , $Q_{50}(\text{median})$, Q_{75} separately for males and females.
- (3 points) The original census cards have been scanned are available online. [This one in particular](#) is quite famous. Why?

4. (25 points) Flint Blood Lead Levels

Lead in the environment is persistent, bio-accumulative, and toxic. Chronic exposure to lead in children is associated with many negative health outcomes even when the Blood Lead Levels (BLLs) are measured as low as 1.0-10.0 $\mu\text{g}/\text{dL}$. An analysis of childhood exposure to lead is described in the article *Blood Lead Levels of Children in Flint, Michigan: 2006-2016* by Gomez et al. (2018) available on myCourses.

- a) (2 points) Summarize the main findings of the study in 280 characters or less.
- b) (5 points) Does Figure 1 do a good job of conveying its message? Explain why or why not.
- c) (6 points) Consider the information presented in Figure 1 and think about the dataset which would have been used to generate the plot. What are the rows and what are the columns? What is the dimension of the dataset? Are the data in tidy format?
- d) (6 points) From the graph, extract the yearly BLL percentages $\geq 5.0 \mu\text{g}/\text{dL}$, in children 5 years and younger, for Flint and Michigan. You may read directly off the graph or try using a [WebPlotDigitizer](#). Calculate the yearly change from baseline (2006) percentages separately for each group. Is there evidence to suggest that the change from baselines are different for Flint vs. Michigan? Support your answer with summary statistics and/or a plot.
- e) (6 points) Figure 2 shows the geometric mean BLL levels over time for children residing within Flint boundaries. Recreate the plot and add the simple linear regression line. Calculate the correlation coefficient r and compare it with the linear regression output. Do the results agree?